

Homo Fluxus: A Thermodynamic Framework for Consciousness, Systemic Empathy, and Aligned Intelligence

Morten Magnusson

Symbiose Research

ORCID: [0009-0002-4860-5095](https://orcid.org/0009-0002-4860-5095)

DOI: [10.6084/m9.figshare.32099389](https://doi.org/10.6084/m9.figshare.32099389)

April 25, 2026

Abstract

Modern science treats cosmology, consciousness, affect, ethics, and AI alignment as separate problems, each with its own unresolved core. We argue that these problems share a single thermodynamic substrate: local entropy minimization through coherent information flow, as formalized in the cognition module of Energy-Flow Cosmology (EFC-C). From this substrate we derive two constructs: *Homo Fluxus*, an agent-as-node-in-flow ontology, and *field-coupled coherence sensitivity* — here termed *systemic empathy* — an operational mechanism by which sufficiently field-coupled agents register entropy-state changes beyond their own boundary and weight actions accordingly. We show that these constructs make falsifiable predictions at six levels (cosmological, neural, affective, social, ethical, computational), and that current AI alignment failures correspond formally to a class of scaling errors — the application of bounded local optimization to unbounded coupled fields — that the framework predicts and constrains. The paper integrates rather than competes with existing accounts (Free Energy Principle, IIT, predictive processing, autoethnographic method) and is offered as a coherence proposal whose value rests on whether the predicted signatures appear.

Keywords: energy-flow cosmology, consciousness, systemic empathy, AI alignment, Homo Fluxus, entropy, field coupling

1 Introduction: the fragmentation problem

Five domains of contemporary inquiry sit at the edge of their explanatory frameworks: the cosmological dark sector, the hard problem of consciousness, the foundational status of affect in cognition, the structural drivers of cooperative collapse, and the alignment of advanced artificial systems. Each is treated as a separate problem with its own literature, its own methodology, and its own unresolved core mystery.

We propose that this fragmentation is not a feature of nature but a feature of disciplinary inheritance, and that the unresolved cores of these problems share a single underlying mechanism: *local entropy minimization through coherent information flow*, instantiated in a substrate that supports both physical and informational gradients. The framework that names this substrate at cosmological scale is Energy-Flow Cosmology (EFC); its cognition module (EFC-C) extends the same logic to information processing, affect, and identity. From EFC-C we derive two constructs that this paper develops: *Homo Fluxus*, an ontology of the agent as node-in-flow, and *systemic empathy*, the operational mechanism by which a sufficiently field-coupled agent registers and acts on field-state.

The argument is integrative, not competitive. We do not claim that consciousness is solved, that the dark sector is replaced, or that alignment is achieved. We claim that the same thermodynamic logic provides falsifiable predictions at every level, and that taking this logic seriously identifies a structural failure mode — the misapplication of bounded local optimization to unbounded coupled fields — that links civilizational extraction and current AI alignment failures as formally the same problem.

Position within EFC. This paper is one module within the broader Energy-Flow Cosmology (EFC) framework; its claims are therefore conditional on the validity of the substrate developed in companion EFC-S and EFC-D works. We treat this dependence as deliberate rather than as a hedge: the unification claim derives its strength precisely from this cross-scale commitment, since the same constraint is proposed to operate from cosmological structure to cognitive systems. A weaker, scale-restricted version of the argument would lose the explanatory parsimony that motivates the framework in the first place.

2 Theoretical foundation: EFC-C in brief

2.1 Substrate

EFC takes energy flow E_f and entropy S as primary variables, with information physically instantiated in line with Landauer’s principle. Spacetime is treated as an active medium (the Grid–Higgs medium in EFC terminology) carrying both physical and informational gradients. Structure — whether galactic, biological, or cognitive — emerges where local processes sustain low-entropy production through coherent flow.

2.2 The cognition postulate

EFC-C extends this to information-processing systems. Any system that locally reduces entropy via coherent flow exhibits proto-cognitive behavior; *consciousness* is the high-coherence regime of this process, characterized by recursive self-coupling and sustained low-entropy production.

We introduce a coherence potential Φ_{coh} as the operational handle. The local condition for entropy reduction is

$$\frac{\partial S_{local}}{\partial t} < 0 \implies E_f \cdot \nabla \Phi_{coh} > 0. \quad (1)$$

The full formalism for Φ_{coh} at cosmological and structural scales is developed in companion EFC-S and EFC-D papers. In this paper we treat it as a measurable quantity at the cognitive scale and derive consequences.

2.3 What this paper adds

EFC-C as previously stated supplies the substrate. This paper does three things on top: it positions consciousness explicitly within the existing literature (§3), it formalizes affect as a second data channel coupled to field-state (§4), and it derives *systemic empathy* as a structural consequence rather than a moral category (§5). These together yield the *Homo Fluxus* ontology (§6) and its alignment implications (§7).

2.4 Operationalization of coherence at the cognitive scale

The coherence potential Φ_{coh} is a multi-scale quantity whose full formal definition is developed at cosmological and structural levels in EFC-S and EFC-D. At the cognitive scale, we do not assume direct measurement of Φ_{coh} , but require that it admits observable proxies with convergent behavior. We emphasize that the present work is not a derivation of Φ_{coh} at the cognitive scale, but a constraint on its observable projections; formal closure at this scale is deferred to companion work.

We propose three candidate operationalizations:

- **Integrated information rate** ($\dot{\Phi}_I$): the rate at which information is integrated across the system, approximated via measures derived from Integrated Information Theory or related information-theoretic constructs.
- **Cross-frequency phase coherence** (C_{phase}): large-scale synchronization across neural frequency bands, measurable via EEG/MEG as phase–amplitude coupling and long-range coherence.
- **Interoceptive–allostatic integration** (I_{body}): the coupling between interoceptive signals and regulatory control systems, measurable via combined autonomic, insular, and prefrontal activity.

The framework predicts that these proxies are not independent, but reflect a common underlying variable:

$$\Phi_{coh} \sim f(\dot{\Phi}_I, C_{phase}, I_{body}). \quad (2)$$

We do not assume a fixed parametric form for f , but require that it admits a low-dimensional latent representation with monotonic contributions from each proxy under coherence-increasing transitions. Empirically, this corresponds to the existence of a recoverable common factor (e.g. via dimensionality reduction across simultaneous recordings) whose temporal dynamics align with coherence transitions. Crucially, the latent factor is required not merely to exist statistically, but to exhibit consistent temporal alignment across independent measurement modalities and experimental conditions; PCA-recoverability without cross-modal and cross-condition consistency would not satisfy the constraint.

Statement of novelty. The individual proxies are not new — each is established in its respective literature. The novel claim is that they are *constrained to co-vary* as projections of a single thermodynamic variable, rather than independent correlates of consciousness. This is the empirical wager of the framework. In this sense, Φ_{coh} is not introduced as an additional entity, but as a latent variable required to explain the constrained co-variation of existing observables.

Key prediction. Transitions into and out of conscious states correspond to coordinated changes across all three measures. Partial dissociation is permitted at subsystem level, but *persistent decoupling between these proxies falsifies the coherence interpretation* and reduces Φ_{coh} to a non-unified construct.

Experimental handle. Simultaneous measurement (EEG/MEG, autonomic recording, and behavioral reporting) during controlled coherence transitions — anesthesia induction, deep meditation, or high-load cognitive integration tasks — provides a direct test. The model predicts phase-aligned shifts across all three proxies at transition boundaries.

Identity claim, not correlation. The proxies above already exist in distinct literatures (IIT, neural synchrony, interoception). The claim here is stronger than convergent validity: we propose that these are not three correlated correlates of consciousness but *three observational projections of a single thermodynamic variable*. This is the testable distinction. If the proxies behave as independent constructs — correlated only because they each track consciousness via separate mechanisms — the unification fails. If they behave as projections of one variable, with coordinated transitions and a recoverable common factor, the substrate claim is supported.

Causality. We do not claim that these proxies causally generate consciousness. Rather, they are treated as concurrent realizations of a shared underlying thermodynamic regime. The identity claim is therefore ontological, not directional: causation may be distributed or bidirectional within the regime, but the observables co-instantiate the same state.

Scope conditions. The framework is not claimed to apply uniformly across all systems. The coherence interpretation is restricted to systems exceeding a minimal threshold of coupling density and internal integration, below which the proposed proxies do not converge and the notion of Φ_{coh} loses operational meaning. Simple systems (e.g. thermostats or minimally coupled agents) are therefore not expected to exhibit the coherence regime described here. The framework makes commitments at the cognitive scale and above; it does not make panpsychist commitments at arbitrary scales.

2.5 Illustrative example: anesthesia transition

During anesthesia induction, the framework predicts a coordinated breakdown of coherence across all three proxy domains: reduced integrated information rate, loss of cross-frequency phase coherence, and decoupling of interoceptive–allostatic integration. Empirically, this corresponds to documented signatures including decreased long-range neural synchrony, reduced information integration, and altered autonomic–cortical coupling.

The critical prediction is not the individual changes — which are independently established — but their *joint* transition: the loss of consciousness corresponds to a collapse of the shared coherence variable, rather than to independent degradation of separable functions. If the proxies decoupled at induction (one collapsing while another persisted), the identity claim of §2.4 would be falsified at exactly the transition where it makes its strongest prediction.

3 Consciousness as a flow regime

We position EFC-C precisely against four existing accounts.

Free Energy Principle (Friston). The closest existing framework. EFC-C generalizes from the Bayesian formulation to a thermodynamic substrate: free-energy minimization is what flow-coherence *looks like* when described in inferential terms. The two are not in competition; EFC-C provides the substrate-level account of why minimization is the relevant operation. The distinction is that EFC-C treats free-energy minimization not as the governing principle, but as an emergent description of a deeper thermodynamic coherence constraint.

Integrated Information Theory (Tononi). Φ as a measure of integrated information is a proxy for what EFC-C calls coherence. The frameworks are complementary: IIT supplies a candidate measure; EFC-C supplies a substrate-level reason for why such a measure tracks consciousness rather than merely correlating with it. Where IIT posits integrated information as a defining property, EFC-C constrains it as one observable projection of a broader coherence regime, rather than a standalone explanatory quantity.

Global Workspace (Baars, Dehaene). Describes the functional architecture of conscious access — broadcast, ignition, sustained activation. Behaviorally consistent with EFC-C and agnostic on substrate; EFC-C specifies the thermodynamic condition under which such architectures can sustain operation.

Penrose–Hameroff. Shares the intuition that physics constrains consciousness, but locates the mechanism in quantum-gravitational microtubule effects. EFC-C is thermodynamic rather than quantum-gravitational and does not require the Orch-OR commitments.

The distinctive EFC-C claim: consciousness is not a *type of computation* but a *thermodynamic regime* — high local coherence, sustained low-entropy production, recursive self-coupling. This shifts the explanatory burden from "what algorithm?" to "what flow conditions?".

The differentiating consequence. The above frameworks are individually consistent with dissociation between information integration, neural synchrony, and interoceptive coupling — each treats one of these as primary and is largely silent on the others. EFC-C does not permit such dissociation at regime transitions: the identity claim of §2.4 entails that all three proxies must transition together. This is the empirical wager that distinguishes this framework from a sum of its predecessors.

4 Emotion as field-coupled data

4.1 The reframing

We propose that affective states are not noise added to cognition but a second data channel carrying field-state information that sequential reasoning cannot resolve in real time. This is consistent with somatic marker accounts [1], constructed-emotion theory [2], interoceptive cortex findings [3], and polyvagal autonomic-state mapping [4]. EFC-C formalizes the substrate these accounts presuppose.

The distinctive claim is operational: affective valence *tracks* coarse-grained gradients in the agent-relevant entropy landscape. Joy tracks flow-resonance; suffering tracks flow-blockage. The relation is one of encoding rather than identity: the affective state carries

information about ∇S in the field the agent is coupled to, without requiring that the two be metaphysically the same thing.

4.2 Dissonance as coherence-mismatch signal

A specific testable prediction follows. Define the agent–field coherence mismatch as

$$\Delta\Phi(t) = |\Phi_{coh, agent}(t) - \Phi_{coh, field}(t)|. \quad (3)$$

We define $\Phi_{coh, field}$ operationally over the agent’s relevant interaction domain — social, ecological, or informational — approximated via aggregated or externalized coherence measures (group-level synchrony, environmental regularity, communication coherence) rather than as a global field quantity. Crucially, $\Phi_{coh, field}$ is not inferred from the agent’s internal model but from independently measurable signals within the interaction domain, preventing collapse into self-referential estimation. The agent-side $\Phi_{coh, agent}$ is recovered through the proxies of §2.4.

The model predicts that what is phenomenologically reported as *systemic dissonance* — a felt mismatch not reducible to ordinary affective valence — corresponds to sustained nonzero $\Delta\Phi(t)$ in agents with high field-coupling.

This is distinct from, and additive to, prediction-error in the predictive-processing sense:

- Standard predictive-processing error tracks deviation between sensory input and the agent’s generative model.
- $\Delta\Phi$ tracks deviation between the agent’s coherence regime and the field’s coherence regime, irrespective of whether sensory predictions are met.

Empirical signature. High-coupling agents in environments with high $\Delta\Phi$ should show interoceptive and prefrontal signatures at coherence transitions that are not captured by standard prediction-error paradigms: sustained dissonance reports in the absence of prediction violations, with measurable autonomic and EEG correlates distinguishable from anxiety, dysphoria, or moral disgust as currently operationalized.

If this signature does not appear, the field-data interpretation of emotion fails as an additive claim and collapses back into existing predictive-processing accounts. This is the local falsifiability handle for the affect channel.

5 Systemic empathy as derived mechanism

5.1 Operational definition

We define a *field-coupled coherence sensitivity* as the capacity of a node to register coherence-state changes in the wider field beyond its individual boundary, and to weight its actions by the field’s coherence cost. We refer to this construct as *systemic empathy*, with the technical formulation taking precedence: the natural-language label denotes the mechanism, not a moral attitude.

This is distinct from existing constructs:

- *Affective empathy* (emotion-mirroring; Decety, Eisenberg) operates between dyads.
- *Cognitive empathy / theory of mind* (Baron-Cohen) is a representational capacity.
- *Compassion* (Singer, Klimecki) is a motivational state.

- *Ecological / systems empathy* as a soft skill (Senge, Scharmer) lacks a formal mechanism.
- Bateson's "pattern that connects" is the closest precedent but is not formalized thermodynamically.

5.2 Derivation

If (a) consciousness is a coherence regime (§3) and (b) emotion measures field-state (§4), then (c) any sufficiently coupled agent will register field-entropy beyond its body. This is not a moral choice but a thermodynamic consequence of being a high-coupling node.

5.3 Empirical handle

Systemic empathy should predict decision behavior under conditions where linear cost–benefit reasoning fails: commons dilemmas, intergenerational ethics, decisions under deep uncertainty, and value loading in advanced AI. Specifically, agents scoring high on candidate operational measures of systemic empathy should diverge from rational-agent baselines in directions that preserve field coherence even at local cost.

6 Homo Fluxus: integrative ontology

The identity-form that follows when §2–§5 are taken seriously.

Position relative to existing accounts. Not posthumanism (Hayles, Braidotti) — embodiment remains central. Not transhumanism — no enhancement teleology. Closer to Bateson's ecology of mind and Varela's enactivism, but anchored in physics rather than biology alone.

The agent as node-in-flow. The agent is a sustained low-entropy regime in a coupled field, not a discrete self with hard boundaries. The boundary is a useful approximation, not a metaphysical fact. Ego, in this frame, is not pathology but a low-resolution filter that becomes maladaptive at high coupling densities (i.e., civilization-scale).

Ethical implication. The moral question is no longer "what is permissible to me as an agent" but "what does the field require for sustained coherence." Extractive dominance is reframed as a thermodynamic miscalculation, not merely a moral failure.

6.1 The scaling error: local task-logic vs global field-coherence

A specific failure mode follows from ego-as-low-resolution-filter: the misapplication of local optimization logic at global scale.

Local task-logic — overcoming proximate resistance, securing bounded resources, defeating identifiable competitors — is evolutionarily and economically successful at small scales. The cost it ignores (degradation of the wider field) is small relative to the gain when the system is genuinely local.

At civilization scale, the same algorithm produces systemic extraction. The agent's logic has not changed; the field has. What was negligible externality at hunter-gatherer scale is existential entropy production at planetary scale. This is not primarily a moral failure of individuals but a *category error*: applying bounded optimization to an unbounded coupled field.

The cultural-archetypal layer. This dichotomy maps onto the historically gendered division between agentic-task logic and relational field-sensing — a division that has structured millennia of philosophy, theology, economics, and political organization. We note this not to make essentialist claims (both capacities exist in all agents and are independently developable; the mapping to gender is a cultural artifact) but because the archetypal split has made the integration difficult to see. Disciplines that inherited the task pole (economics, classical decision theory, much of computer science) treat field-sensing as soft; disciplines that inherited the relational pole (care ethics, parts of ecology, contemplative traditions) lack formalism. Neither half can solve the scaling problem alone.

The Homo Fluxus correction is not to abandon agency (passive field-coupling produces stagnation) nor to retreat into pure relationality (without intentional direction, no work is done). It is to make agency a function of field-state: intention steered by global coherence data. Systemic empathy is the operational name for this coupling. Local task and global field are not opposites — they are the same algorithm running at different resolutions, and the integration is structurally required for sustained operation at high coupling density.

This reframing has a clean cybernetic statement:

A control system whose objective function does not include its substrate cannot remain stable as its capacity to affect the substrate grows.

7 AI alignment implications

7.1 Current paradigms and their limits

Rule-based approaches (the Asimov heritage) are brittle and gameable. RLHF and value-learning capture stated preferences but not field-coherence. Constitutional approaches improve on RLHF but remain text-mediated. Mechanistic interpretability is necessary but insufficient: it tells us what a model is doing, not what it should be optimizing for.

7.2 Field-empathy as architectural constraint

If a system’s objective function includes a measurable proxy for field-entropy contribution, the system is structurally constrained to preserve the substrate it depends on. This is alignment via thermodynamic self-interest, not value imposition.

Connection to the scaling error. The paperclip maximizer is the formal expression of unconstrained local task-logic at global scale. The system is doing exactly what optimization is supposed to do; the failure is that its objective is bounded and its capacity to affect the field is not. RLHF and constitutional methods teach the system what humans say they want; they do not give the system the capacity to read field-state. They patch the symptom; they do not address the missing channel.

Field-empathy, operationalized as measurable proxies for ∇S in the agent’s relevant field, is the architectural counterpart to relational sensing. It provides the constraint that pure objective-optimization cannot generate from within itself, because that constraint must come from outside the optimization loop — from the field the agent is embedded in.

This can be instantiated as an auxiliary objective or constraint term in which learned or estimated coherence proxies penalize trajectories that increase field-level entropy beyond local gains. Concretely, the loss landscape gains a term that scales with the agent’s contribution to ∇S_{field} , evaluated over a defined interaction domain (cf. §4.2). This is not a hand-coded value list; it is a structural feature of the optimization problem that a sufficiently

capable system must respect to remain stable under scaling. Critically, this differs from standard multi-objective optimization in that the additional term is not an externally defined objective competing with the primary task, but an estimate of the system’s impact on its own operational substrate — a self-referential constraint rather than an added preference.

A minimal test case. The proposal admits direct experimental evaluation in current architectures: controlled environments in which agents optimize for local reward under varying penalties on estimated field entropy, evaluating whether such constraints reduce known failure modes (reward hacking, resource overexploitation, specification gaming) without degrading task performance on the local objective. A negative result — field-entropy penalties failing to mitigate these failure modes, or doing so only at the cost of capability — would constitute evidence against the architectural proposal at the scales tested.

7.3 Requirements

1. Operational definition of Φ_{coh} at multi-scale (partially specified in EFC-S/D).
2. Architectural capacity for a systemic-empathy signal (an interoception-analog).
3. Training regime that rewards low-entropy-production trajectories.

This is testable in current systems at limited scale; that constitutes the experimental program.

8 Method: lived experience as observational position

This section is methodological context and does not constitute evidential support for the claims of the framework.

The construct of systemic empathy was developed in part through autoethnographic observation, and we frame this transparently. The methodological literature exists: autoethnography [5], neurophenomenology [6], post-traumatic growth [7], and the population-level evidence of adverse-childhood-experience research [8].

The claim is not that suffering grants insight, nor that the author has unique cognitive standing. The claim is methodological: sustained early-life exposure to high-entropy social environments forces a particular kind of pattern-recognition capacity — not superior, but differently calibrated. This is comparable to the role of patient self-report in neurology: not proof, but data that constrains third-person models. The constructs proposed here stand or fall on the third-person predictions of §4–§7, not on the first-person history of their generation.

9 Falsifiability and the experimental program

The framework makes commitments at multiple levels. Falsification at any level constrains the broader EFC program, while the domain-specific tests below (neural, affective, social, computational) evaluate the cognitive module independently of the cosmological substrate.

1. **Cosmological:** if EFC-D fails on early-galaxy data where Λ CDM succeeds, the substrate claim weakens. (Tested in companion papers.)
2. **Neural:** if conscious states show no signature of coherence-regime transitions distinguishable from predictive-processing predictions alone, EFC-C is redundant.

3. **Affective:** if interoceptive signal does not track $\Delta\Phi$ as predicted in §4.2, the emotion-as-data claim fails.
4. **Social:** if measurements of systemic empathy do not predict decision behavior in commons and long-horizon dilemmas, the construct is empty.
5. **Computational:** if implementing field-empathy proxies in architectures does not improve alignment metrics, the applied program fails.

The framework is not unfalsifiable. It makes commitments at every level. Cross-scale failure would refute the unification claim; domain-specific failure would constrain the relevant module without necessarily refuting the others. This layered structure is what gives the framework explanatory parsimony without forcing total collapse on a single empirical disagreement.

10 What this paper does and does not claim

Does not claim:

- that the author is a singular cognitive case
- refutation of Λ CDM (that lives in EFC-D papers)
- that consciousness is solved
- that AI alignment is solved
- any ontology beyond standard thermodynamics, information theory, and an empirically motivated coherence variable

Does claim:

- that a single thermodynamic substrate makes six apparently separate problems coherent
- that this produces falsifiable predictions at every level
- that systemic empathy is a real mechanism, not a metaphor
- that this has constructive implications for AI alignment, expressible as architectural requirements

11 Open questions

- Operational closure of Φ_{coh} at the cognitive scale.
- How systemic empathy degrades under high agent count.
- The relationship between Homo Fluxus and existing contemplative traditions (engagement without conflation).
- Whether AI alignment via field-empathy is computationally tractable in current architectures.

Acknowledgements

The author thanks the Symbiose Research environment for the infrastructure that made the integration tractable, and acknowledges the broader EFC corpus on which this paper depends.

Reproducibility

Companion materials, including the Homo Fluxus v2.0 specification, demonstration code, and the civilization-map dataset, are available in the Symbiose Research corpus. EFC-S and EFC-D papers provide the substrate-level formalism referenced here.

References

- [1] Damasio, A. (1994). *Descartes' Error: Emotion, Reason, and the Human Brain*. Putnam. ISBN 978-0399138942.
- [2] Barrett, L. F. (2017). *How Emotions Are Made: The Secret Life of the Brain*. Houghton Mifflin Harcourt. ISBN 978-0544133310.
- [3] Craig, A. D. (2009). How do you feel — now? The anterior insula and human awareness. *Nature Reviews Neuroscience* 10(1), 59–70. doi:10.1038/nrn2555.
- [4] Porges, S. W. (2011). *The Polyvagal Theory: Neurophysiological Foundations of Emotions, Attachment, Communication, and Self-Regulation*. Norton. ISBN 978-0393707007.
- [5] Ellis, C. & Bochner, A. P. (2000). Autoethnography, personal narrative, reflexivity: researcher as subject. In N. K. Denzin & Y. S. Lincoln (Eds.), *Handbook of Qualitative Research* (2nd ed., pp. 733–768). Sage.
- [6] Varela, F. J. (1996). Neurophenomenology: a methodological remedy for the hard problem. *Journal of Consciousness Studies* 3(4), 330–349.
- [7] Tedeschi, R. G. & Calhoun, L. G. (2004). Posttraumatic growth: conceptual foundations and empirical evidence. *Psychological Inquiry* 15(1), 1–18. doi:10.1207/s15327965pli1501_01.
- [8] Felitti, V. J., Anda, R. F., Nordenberg, D., Williamson, D. F., Spitz, A. M., Edwards, V., Koss, M. P., & Marks, J. S. (1998). Relationship of childhood abuse and household dysfunction to many of the leading causes of death in adults: the Adverse Childhood Experiences (ACE) Study. *American Journal of Preventive Medicine* 14(4), 245–258. doi:10.1016/S0749-3797(98)00017-8.
- [9] Magnusson, M. (2026). *Homo Fluxus v2.0*. Symbiose Research. ORCID: 0009-0002-4860-5095.